

QTM 220 Final

AUTHOR

Veronica Vargas

```
## HW 1-4
```

```
library(tidyverse)
```

Warning: package 'tidyverse' was built under R version 4.3.3

— Attaching core tidyverse packages — tidyverse 2.0.0 —

✓ dplyr 1.1.3 ✓ readr 2.1.4

✓ forcats 1.0.0 ✓ stringr 1.5.0

✓ ggplot2 3.4.3 ✓ tibble 3.2.1

✓ lubridate 1.9.2 ✓ tidyr 1.3.0

✓ purrr 1.0.2

— Conflicts — tidyverse_conflicts() —

✗ dplyr::filter() masks stats::filter()

✗ dplyr::lag() masks stats::lag()

ℹ Use the conflicted package (<<http://conflicted.r-lib.org/>>) to force all conflicts to become errors

```
library(ggplot2)
```

```
library(dplyr)
```

```
## HW 6 - Cross Validation
```

```
library(mosaic)
```

Warning: package 'mosaic' was built under R version 4.3.3

Registered S3 method overwritten by 'mosaic':

method	from
fortify.SpatialPolygonsDataFrame	ggplot2

The 'mosaic' package masks several functions from core packages in order to add additional features. The original behavior of these functions should not be affected by this.

Attaching package: 'mosaic'

The following object is masked from 'package:Matrix':

mean

The following objects are masked from 'package:dplyr':

count, do, tally

The following object is masked from 'package:purrr':

cross

The following object is masked from 'package:ggplot2':

stat

The following objects are masked from 'package:stats':

binom.test, cor, cor.test, cov, fivenum, IQR, median, prop.test,
quantile, sd, t.test, var

The following objects are masked from 'package:base':

max, mean, min, prod, range, sample, sum

```
library(mosaicData)  
library(leaps)
```

Warning: package 'leaps' was built under R version 4.3.3

```
library(caret)
```

Warning: package 'caret' was built under R version 4.3.3

Attaching package: 'caret'

The following object is masked from 'package:mosaic':

dotPlot

The following object is masked from 'package:purrr':

lift

```
library(ISLR2)
```

Warning: package 'ISLR2' was built under R version 4.3.3

```
## HW 7 - Confidence Intervals for Regression Modeling  
library(datasets)  
library(boot)
```

Warning: package 'boot' was built under R version 4.3.3

Attaching package: 'boot'

The following object is masked from 'package:mosaic':

logit

The following object is masked from 'package:lattice':

melanoma

```
library(lmtest) # sandwich package
```

Warning: package 'lmtest' was built under R version 4.3.3

Loading required package: zoo

Warning: package 'zoo' was built under R version 4.3.3

Attaching package: 'zoo'

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

```
## HW 8 - Diagnostic Plots
```

```
library(experimentr)
```

```
library(datasets)
```

```
pokemon.data <- read.csv("C:/Users/13015/OneDrive - Emory University/Documents/Fall 2024/QTM 220/1")
```

```
pokemon.data <- pokemon.data %>%  
  filter(Generation %in% c(1, 2, 3, 4, 5, 6, 7))
```

```
head(pokemon.data)
```

	Name	Generation	Type1	HP	Attack	Defense	SP_Attack	SP_Defense	Speed
1	Bulbasaur	1	Grass	45	49	49	65	65	45
2	Ivysaur	1	Grass	60	62	63	80	80	60
3	Venusaur	1	Grass	80	82	83	100	100	80
4	Mega Venusaur	1	Grass	80	100	123	122	120	80
5	Charmander	1	Fire	39	52	43	60	50	65
6	Charmeleon	1	Fire	58	64	58	80	65	80
	predicted_speed								
1				56.23955					
2				62.26298					
3				70.29423					
4				79.12860					
5				67.27336					
6				71.99359					

```
summary(pokemon.data)
```

Name	Generation	Type1	HP
Length:940	Min. :1.000	Length:940	Min. : 1.00
Class :character	1st Qu.:2.000	Class :character	1st Qu.: 50.00
Mode :character	Median :4.000	Mode :character	Median : 66.00
	Mean :3.818		Mean : 69.55
	3rd Qu.:5.000		3rd Qu.: 80.00
	Max. :7.000		Max. :255.00
Attack	Defense	SP_Attack	SP_Defense
Min. : 5.00	Min. : 5.00	Min. : 10.00	Min. : 20.00
1st Qu.: 55.00	1st Qu.: 50.00	1st Qu.: 50.00	1st Qu.: 50.00
Median : 75.00	Median : 70.00	Median : 65.00	Median : 70.00
Mean : 80.05	Mean : 74.38	Mean : 73.28	Mean : 72.12
3rd Qu.:100.00	3rd Qu.: 90.00	3rd Qu.: 95.00	3rd Qu.: 90.00
Max. :190.00	Max. :230.00	Max. :194.00	Max. :230.00
Speed	predicted_speed		
Min. : 5.0	Min. : 24.54		
1st Qu.: 45.0	1st Qu.: 58.21		
Median : 65.0	Median : 66.86		
Mean : 68.8	Mean : 68.80		
3rd Qu.: 90.0	3rd Qu.: 78.84		
Max. :180.0	Max. :131.00		

Exercise #1

(a) Scatter Plot #1-2

```
mod.simple <- lm(Speed ~ SP_Defense, data = pokemon.data)
summary(mod.simple)
```

Call:

```
lm(formula = Speed ~ SP_Defense, data = pokemon.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-101.80	-21.89	-1.07	20.89	106.90

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	51.44342	2.62866	19.570	<2e-16 ***
SP_Defense	0.24067	0.03404	7.071	3e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 28.84 on 938 degrees of freedom

Multiple R-squared: 0.05061, Adjusted R-squared: 0.04959


F-statistic: 50 on 1 and 938 DF, p-value: 3.004e-12

```
pokemon.data$predicted_score_simple <- predict(mod.simple)

ggplot(pokemon.data, aes(x = SP_Defense, y = Speed)) +
  geom_point(aes(x = SP_Defense, y = Speed),
             alpha = 0.3) +
  geom_line(aes(y = predicted_score_simple), color = "orange", size = 1) +
  labs(
    title = "SP Defense vs. Speed",
    x = "SP Defense",
    y = "Speed") +
  theme_minimal()
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.

• Please use `linewidth` instead.



```
mod.simple <- lm(Speed ~ SP_Attack, data = pokemon.data)
summary(mod.simple)
```

Call:

```
lm(formula = Speed ~ SP_Attack, data = pokemon.data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-78.645	-18.685	-0.354	17.970	102.161

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	38.30555	2.07103	18.50	<2e-16 ***
SP_Attack	0.41614	0.02575	16.16	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 26.18 on 938 degrees of freedom


Multiple R-squared: 0.2178, Adjusted R-squared: 0.217

F-statistic: 261.2 on 1 and 938 DF, p-value: < 2.2e-16

```
pokemon.data$predicted_score_simple <- predict(mod.simple)

ggplot(pokemon.data, aes(x = SP_Attack, y = Speed)) +
  geom_point(aes(x = SP_Attack, y = Speed),
             alpha = 0.3) +
  geom_line(aes(y = predicted_score_simple), color = "orange", size = 1) +
  labs(
    title = "SP Attack vs. Speed",
    x = "SP Attack",
    y = "Speed") +
  theme_minimal()
```

SP Attack vs. Speed



SP Attack is more highly correlated with Speed. When we're looking at the coefficients for each simple regression model, the coefficient between SP Attack and Speed is 0.41614, which is larger than the coefficient between SP Defense and Speed which is 0.24067.

(b) Quantile w/ Bootstrapped Estimated Sampling Distribution

```
quantile(pokemon.data$SP_Attack, 0.85)
```

85%
108.15

```
n <- 10000
df <- rep(NA, n)


for(i in 1:n){
  sample <- sample(pokemon.data$SP_Attack, 940, replace = T)

  df[i] <- quantile(sample, 0.85)
}
```

```
ggplot(data = data.frame(df = df), aes(x = df)) +
  geom_histogram(fill = "cyan4", alpha = 0.5, Type1 = "identity") +
  geom_vline(xintercept = mean(df), linetype="dashed",
             color = "coral", linewidth=1) +
  geom_vline(xintercept = quantile(df, 0.025), linetype = 'dotted',
             color = "darkorchid", linewidth = 1) +
  geom_vline(xintercept = quantile(df, 0.975), linetype = "dotted",
             color = "darkorchid", linewidth=1) +
  theme_minimal()
```

Warning in geom_histogram(fill = "cyan4", alpha = 0.5, Type1 = "identity"):
Ignoring unknown parameters: `Type1`

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
lower.bound <- quantile(df, 0.05)
upper.bound <- quantile(df, 0.95)

print(paste0("The Bootstrapped 90% CI is {", lower.bound, ", ", upper.bound, "}"))
```

```
[1] "The Bootstrapped 90% CI is {105, 110}"
```


If we were to repeat this experiment under the same conditions with a sufficiently large sample size, the quantile 85 would fall somewhere within the estimated interval for about 95 out of 100 trials. In this experiment, expected value of the quantile 85 for our estimator is predicted to be somewhere in between 105 and 110.

(c) Not-Necessarily Parallel Lines Model #1

```
mod.interaction1 <- lm(Speed ~ SP_Attack + Type1 + SP_Attack*Type1, data = pokemon.data)
summary(mod.interaction1)
```

Call:

```
lm(formula = Speed ~ SP_Attack + Type1 + SP_Attack * Type1, data = pokemon.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-68.148	-16.774	-1.511	16.212	104.246

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.263e+01	6.069e+00	5.377	9.64e-08	***
SP_Attack	5.415e-01	9.615e-02	5.632	2.38e-08	***
Type1Dark	1.505e+01	1.195e+01	1.259	0.20851	
Type1Dragon	1.163e+01	1.212e+01	0.960	0.33754	
Type1Electric	4.099e+01	1.269e+01	3.229	0.00129	**
Type1Fairy	-3.571e+01	1.936e+01	-1.845	0.06539	.
Type1Fighting	3.938e+00	1.162e+01	0.339	0.73485	
Type1Fire	2.048e+01	1.199e+01	1.708	0.08799	.
Type1Flying	-6.756e+00	4.191e+01	-0.161	0.87199	
Type1Ghost	-1.253e+01	1.306e+01	-0.960	0.33733	
Type1Grass	-2.494e+00	1.027e+01	-0.243	0.80824	
Type1Ground	1.082e+01	1.150e+01	0.941	0.34706	
Type1Ice	-3.505e+00	1.417e+01	-0.247	0.80473	
Type1Normal	2.035e+01	8.490e+00	2.397	0.01674	*
Type1Poison	-8.686e+00	1.404e+01	-0.619	0.53617	
Type1Psychic	-3.401e+00	1.037e+01	-0.328	0.74298	
Type1Rock	-1.519e+01	1.017e+01	-1.494	0.13558	
Type1Steel	-1.281e+01	1.231e+01	-1.041	0.29828	
Type1Water	7.241e+00	8.516e+00	0.850	0.39537	
SP_Attack:Type1Dark	-1.410e-01	1.618e-01	-0.872	0.38365	
SP_Attack:Type1Dragon	-1.329e-01	1.410e-01	-0.943	0.34606	
SP_Attack:Type1Electric	-3.877e-01	1.504e-01	-2.577	0.01012	*
SP_Attack:Type1Fairy	1.512e-01	2.410e-01	0.627	0.53061	
SP_Attack:Type1Fighting	6.202e-02	1.926e-01	0.322	0.74749	
SP_Attack:Type1Fire	-3.055e-01	1.463e-01	-2.088	0.03708	*
SP_Attack:Type1Flying	2.715e-01	4.300e-01	0.631	0.52799	
SP_Attack:Type1Ghost	2.525e-04	1.663e-01	0.002	0.99879	
SP_Attack:Type1Grass	-1.399e-01	1.405e-01	-0.996	0.31970	

```

SP_Attack:Type1Ground -1.595e-01 1.860e-01 -0.858 0.39137
SP_Attack:Type1Ice -5.483e-02 1.903e-01 -0.288 0.77331
SP_Attack:Type1Normal -2.340e-01 1.352e-01 -1.731 0.08386 .
SP_Attack:Type1Poison 1.007e-01 2.119e-01 0.475 0.63485
SP_Attack:Type1Psychic -1.913e-02 1.239e-01 -0.154 0.87732
SP_Attack:Type1Rock 1.682e-01 1.482e-01 1.135 0.25650
SP_Attack:Type1Steel -5.255e-02 1.644e-01 -0.320 0.74938
SP_Attack:Type1Water -2.081e-01 1.206e-01 -1.725 0.08481 .

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 25.24 on 904 degrees of freedom

Multiple R-squared: 0.2989, Adjusted R-squared: 0.2718

F-statistic: 11.01 on 35 and 904 DF, p-value: < 2.2e-16

According to the summary, there are least four predictors that is in a relationship with the response variable. These predictor variables are SP_Attack, Type1Electric, Type1Normal, and the SP_Attack*Type1Fire interaction variable. These are distinguishable due to their reported p-value being less than 0.05.

(d) Not-Necessarily Parallel Lines Model #2

```

mod.interaction2 <- lm(Speed ~ SP_Defense + Type1 + SP_Defense*Type1, data = pokemon.data)
summary(mod.interaction2)

```

Call:

```

lm(formula = Speed ~ SP_Defense + Type1 + SP_Defense * Type1,
    data = pokemon.data)

```

Residuals:

Min	1Q	Median	3Q	Max
-69.821	-19.245	-1.964	18.228	98.030

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	58.651220	7.167887	8.182	9.39e-16 ***
SP_Defense	0.066373	0.101193	0.656	0.5120
Type1Dark	6.962557	15.564572	0.447	0.6547
Type1Dragon	-18.942497	16.394883	-1.155	0.2482
Type1Electric	-2.255614	16.051501	-0.141	0.8883
Type1Fairy	-54.511405	21.984922	-2.479	0.0133 *
Type1Fighting	-30.584150	17.369807	-1.761	0.0786 .
Type1Fire	-4.233655	14.373727	-0.295	0.7684
Type1Flying	-58.634271	54.771031	-1.071	0.2847
Type1Ghost	-5.308765	16.394816	-0.324	0.7462
Type1Grass	-24.346870	12.507340	-1.947	0.0519 .
Type1Ground	-23.844350	16.652060	-1.432	0.1525
Type1Ice	-16.207604	14.326395	-1.131	0.2582
Type1Normal	-0.987178	10.156964	-0.097	0.9226

Type1Poison	0.573155	15.898562	0.036	0.9712
Type1Psychic	-3.158566	12.599191	-0.251	0.8021
Type1Rock	1.713205	11.794652	0.145	0.8845
Type1Steel	-32.574608	16.487137	-1.976	0.0485 *
Type1Water	-2.528245	9.697600	-0.261	0.7944
SP_Defense:Type1Dark	0.096684	0.217902	0.444	0.6574
SP_Defense:Type1Dragon	0.421686	0.189263	2.228	0.0261 *
SP_Defense:Type1Electric	0.356461	0.212465	1.678	0.0937 .
SP_Defense:Type1Fairy	0.482971	0.250402	1.929	0.0541 .
SP_Defense:Type1Fighting	0.556814	0.252428	2.206	0.0276 *
SP_Defense:Type1Fire	0.206953	0.194017	1.067	0.2864
SP_Defense:Type1Flying	1.347186	0.731037	1.843	0.0657 .
SP_Defense:Type1Ghost	0.062274	0.209246	0.298	0.7661
SP_Defense:Type1Grass	0.308848	0.171809	1.798	0.0726 .
SP_Defense:Type1Ground	0.407313	0.248950	1.636	0.1022
SP_Defense:Type1Ice	0.226365	0.181028	1.250	0.2115
SP_Defense:Type1Normal	0.135476	0.145570	0.931	0.3523
SP_Defense:Type1Poison	0.013611	0.226189	0.060	0.9520
SP_Defense:Type1Psychic	0.229739	0.150420	1.527	0.1270
SP_Defense:Type1Rock	-0.008551	0.155864	-0.055	0.9563
SP_Defense:Type1Steel	0.295343	0.199762	1.478	0.1396
SP_Defense:Type1Water	0.058492	0.130929	0.447	0.6552

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 27.81 on 904 degrees of freedom

Multiple R-squared: 0.1493, Adjusted R-squared: 0.1164

F-statistic: 4.534 on 35 and 904 DF, p-value: 5.464e-16

The multiple R-squared is 0.1493, meaning that 0.1493 of the variance can be explained by the relationships in the model. Since this number is very low, it's likely that this combination of predictor variables in this multivariate regression do not have much to do with the response variable.

(e) LOOCV

```
rss_summary <- function(data, lev = NULL, model = NULL) {
  residuals <- data$obs - data$pred
  rss <- sum(residuals^2)
  rmse <- sqrt(mean(residuals^2))
  return(c(RMSE = rmse, RSS = rss))
}
```

```
train_control_loocv <- trainControl(
  method = "LOOCV",
  summaryFunction = rss_summary,
  savePredictions = "all",
  classProbs = FALSE,
  allowParallel = FALSE
)
```

```
# Train Model A: Model #1
set.seed(123)
model_A_caret_loocv <- train(
  Speed ~ SP_Attack + Type1 + SP_Attack*Type1,
  data = pokemon.data,
  method = "lm",
  trControl = train_control_loocv,
  metric = "RMSE"
)

# Train Model B: Model #2
set.seed(123)
model_B_caret_loocv <- train(
  Speed ~ SP_Defense + Type1 + SP_Defense*Type1,
  data = pokemon.data,
  method = "lm",
  trControl = train_control_loocv,
  metric = "RMSE"
)


model_A_caret_loocv$results
```

```

pred_data$predicted_speed <- predict(mod.interaction2, newdata = pred_data)

ggplot() +
  geom_point(data = pokemon.data, aes(x = SP_Defense, y = Speed, color = Type1),
            alpha = 0.3) +
  geom_line(data = pred_data, aes(x = SP_Defense, y = predicted_speed,
                                color = Type1), size = 1) +
  labs(
    title = "Sample Non-Parallel Lines Predictive Model",
    x = "Speed",
    y = "SP Defense",
    color = "Type"
  ) +
  theme_minimal()

```



```
coef(mod.interaction2)
```

(Intercept)	SP_Defense	Type1Dark
58.65121971	0.06637337	6.96255726
Type1Dragon	Type1Electric	Type1Fairy
-18.94249740	-2.25561409	-54.51140518


Type1Fighting	Type1Fire	Type1Flying
-30.58414998	-4.23365458	-58.63427056
Type1Ghost	Type1Grass	Type1Ground
-5.30876527	-24.34687003	-23.84434976
Type1Ice	Type1Normal	Type1Poison
-16.20760418	-0.98717761	0.57315475
Type1Psychic	Type1Rock	Type1Steel
-3.15856633	1.71320463	-32.57460791
Type1Water	SP_Defense:Type1Dark	SP_Defense:Type1Dragon
-2.52824537	0.09668430	0.42168623
SP_Defense:Type1Electric	SP_Defense:Type1Fairy	SP_Defense:Type1Fighting
0.35646129	0.48297072	0.55681441
SP_Defense:Type1Fire	SP_Defense:Type1Flying	SP_Defense:Type1Ghost
0.20695298	1.34718595	0.06227394
SP_Defense:Type1Grass	SP_Defense:Type1Ground	SP_Defense:Type1Ice
0.30884774	0.40731319	0.22636459
SP_Defense:Type1Normal	SP_Defense:Type1Poison	SP_Defense:Type1Psychic
0.13547610	0.01361093	0.22973934
SP_Defense:Type1Rock	SP_Defense:Type1Steel	SP_Defense:Type1Water
-0.00855145	0.29534327	0.05849221

Fairy Type Equation

$$\text{Speed} = 58.65 + 0.066 \times \text{SP_Defense} + 0.483 \times \text{SP_Defense:Type1Fairy}$$


(g) Diagnostic Plots

```
mpplot(mod.interaction1, which = 2)
```



```
hist(mod.interaction1$residuals, prob = TRUE, breaks = 20, col = "lightcoral", main = "Residual Histogram")

grid = sort(mod.interaction1$residuals)
lines(grid,
      dnorm(grid,
            mean = mean(mod.interaction1$residuals),
            sd = sd(mod.interaction1$residuals)),
      col = 'maroon', lwd = 2)
```



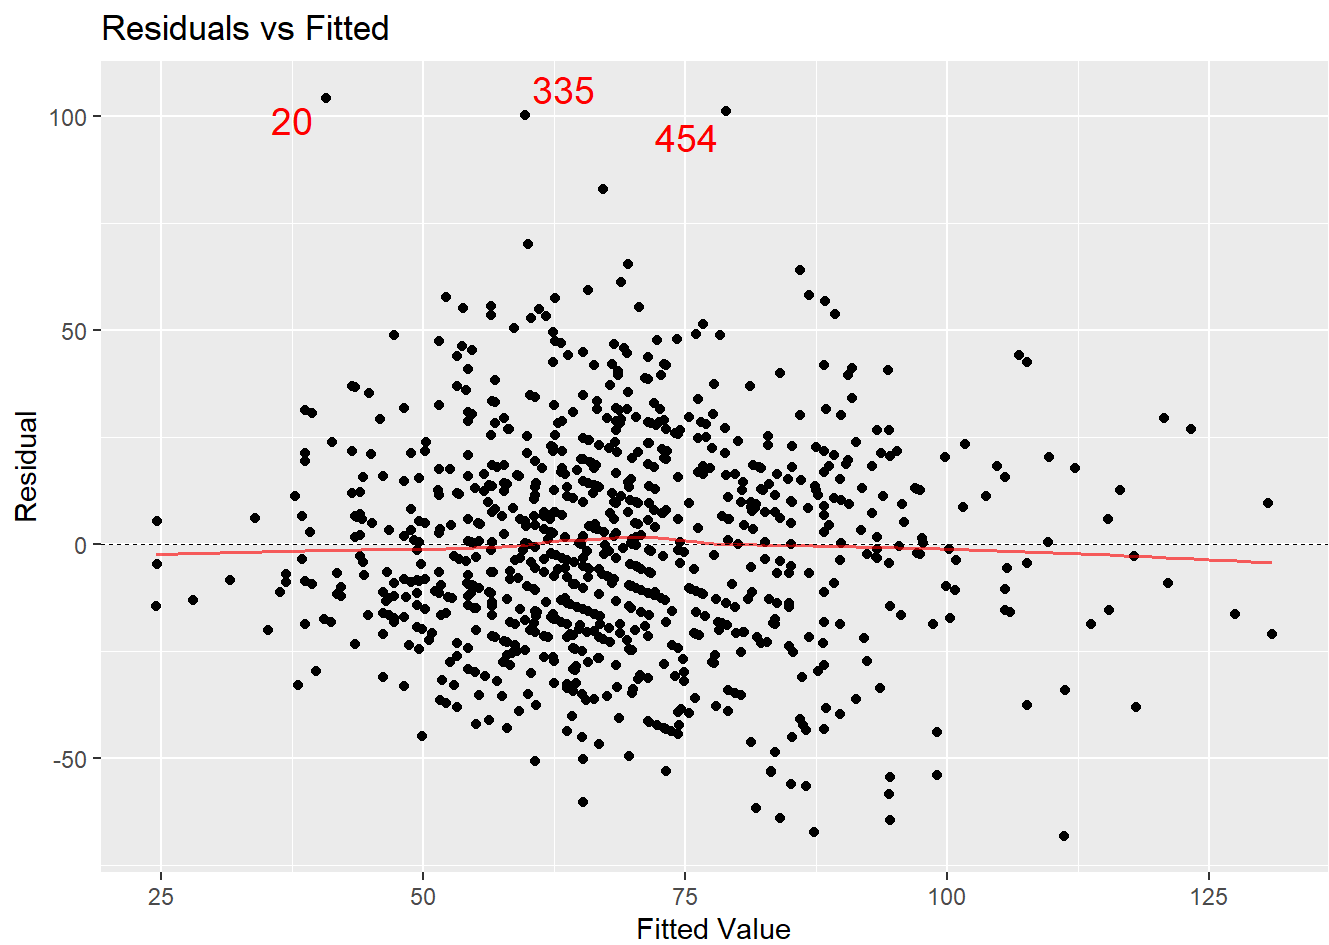
The points follow the line in the Q-Q plot and the histogram follows the density line, meaning that this data meets the assumption of normally distributed residuals.

(h) Mean Zero & Homoscedasticity

```
mpplot(mod.interaction1, which = c(1, 3))
```


```
[[1]]
```

```
`geom_smooth()` using formula = 'y ~ x'
```

```
[[2]]
```

```
`geom_smooth()` using formula = 'y ~ x'
```



While the mean line is around zero, the residuals are scattered all over the place (the residuals are all different), meaning that while this data meets the assumption of mean zero, it does not meet the assumption of homoscedasticity.

(i) 95% Confidence Interval (CI)

```
mod.simple <- lm(Speed ~ SP_Defense, data = pokemon.data)
```

```
coef_hp <- coef(mod.simple)["SP_Defense"]
se_hp <- summary(mod.simple)$coefficients["SP_Defense", "Std. Error"]
```

```
df <- mod.simple$df.residual
t_critical <- qt(1 - 0.05 / 2, df)
```

```
lower_bound <- coef_hp - t_critical * se_hp
upper_bound <- coef_hp + t_critical * se_hp
```

```
lower_bound
```

```
SP_Defense
0.1738758
```

```
upper_bound
```

```
SP_Defense
0.307469
```

```
boot_fn <- function(data, indices) {
  d <- data[indices, ]
  fit <- lm(Speed ~ SP_Defense, data = d)
  return(coef(fit))
}
```

```
set.seed(123)
boot_results <- boot(data = pokemon.data, statistic = boot_fn, R = 10000)
```

```
boot_results
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = pokemon.data, statistic = boot_fn, R = 10000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	51.4434225	-0.0204071449	2.93565814
t2*	0.2406724	0.0004347145	0.04029732

Considering that this model does not meet the assumption of homoscedasticity but does meet the assumptions of mean zero and normality, it would be okay to use the R summary to calculate the 95% CI. The bootstrapped method would be okay as well, but it might be better to use R summary since we are meeting multiple assumptions rather than none.

(j) New Estimator

```
pokemon.subsample <- pokemon.data %>%
  filter ( Type1 %in% c("Dragon", "Bug")) %>%
  select (Name, Generation, Type1, HP )

head(pokemon.subsample)
```

	Name	Generation	Type1	HP
1	Caterpie	1	Bug	45
2	Metapod	1	Bug	50
3	Butterfree	1	Bug	60

4	Weedle	1	Bug	40
5	Kakuna	1	Bug	45
6	Beedrill	1	Bug	65

```
summary(pokemon.subsample)
```

Name	Generation	Type1	HP
Length:117	Min. :1.000	Length:117	Min. : 1.0
Class :character	1st Qu.:3.000	Class :character	1st Qu.: 50.0
Mode :character	Median :4.000	Mode :character	Median : 65.0
	Mean :3.846		Mean : 65.6
	3rd Qu.:5.000		3rd Qu.: 76.0
	Max. :7.000		Max. :216.0

```
## raw difference in mean HP
dragon <- pokemon.subsample[pokemon.subsample$Type1 == "Dragon",]
bug <- pokemon.subsample[pokemon.subsample$Type1 == "Bug",]

mean(dragon$HP)
```

```
[1] 84.08108
```

```
mean(bug$HP)
```

```
[1] 57.05
```

```
mean(dragon$HP) - mean(bug$HP)
```

```
[1] 27.03108
```

```
## difference in mean HP with a focus on bug pokemon

bug <- bug %>%
  group_by(Generation) %>%
  summarise(avg_HP_bug = mean(HP, na.rm = TRUE), n_bug = n()) %>%
  ungroup()

dragon <- dragon %>%
  group_by(Generation) %>%
  summarise(avg_HP_dragon = mean(HP, na.rm = TRUE), n_dragon = n()) %>%
  ungroup()

df <- full_join(bug, dragon, by = "Generation")

df <- df %>%
  mutate(mean_diff = coalesce(avg_HP_bug, 0) - coalesce(avg_HP_dragon, 0))

(1/sum(df$n_bug)) * sum(df$n_bug * df$mean_diff)
```

```
[1] -10.47708
```

```
mod.simple <- lm(HP ~ Type1, data = pokemon.subsample)
summary(mod.simple)
```

Call:

```
lm(formula = HP ~ Type1, data = pokemon.subsample)
```

Residuals:

Min	1Q	Median	3Q	Max
-56.05	-16.08	0.95	12.95	131.92

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	57.050	2.594	21.996	< 2e-16 ***
Type1Dragon	27.031	4.612	5.861	4.48e-08 ***

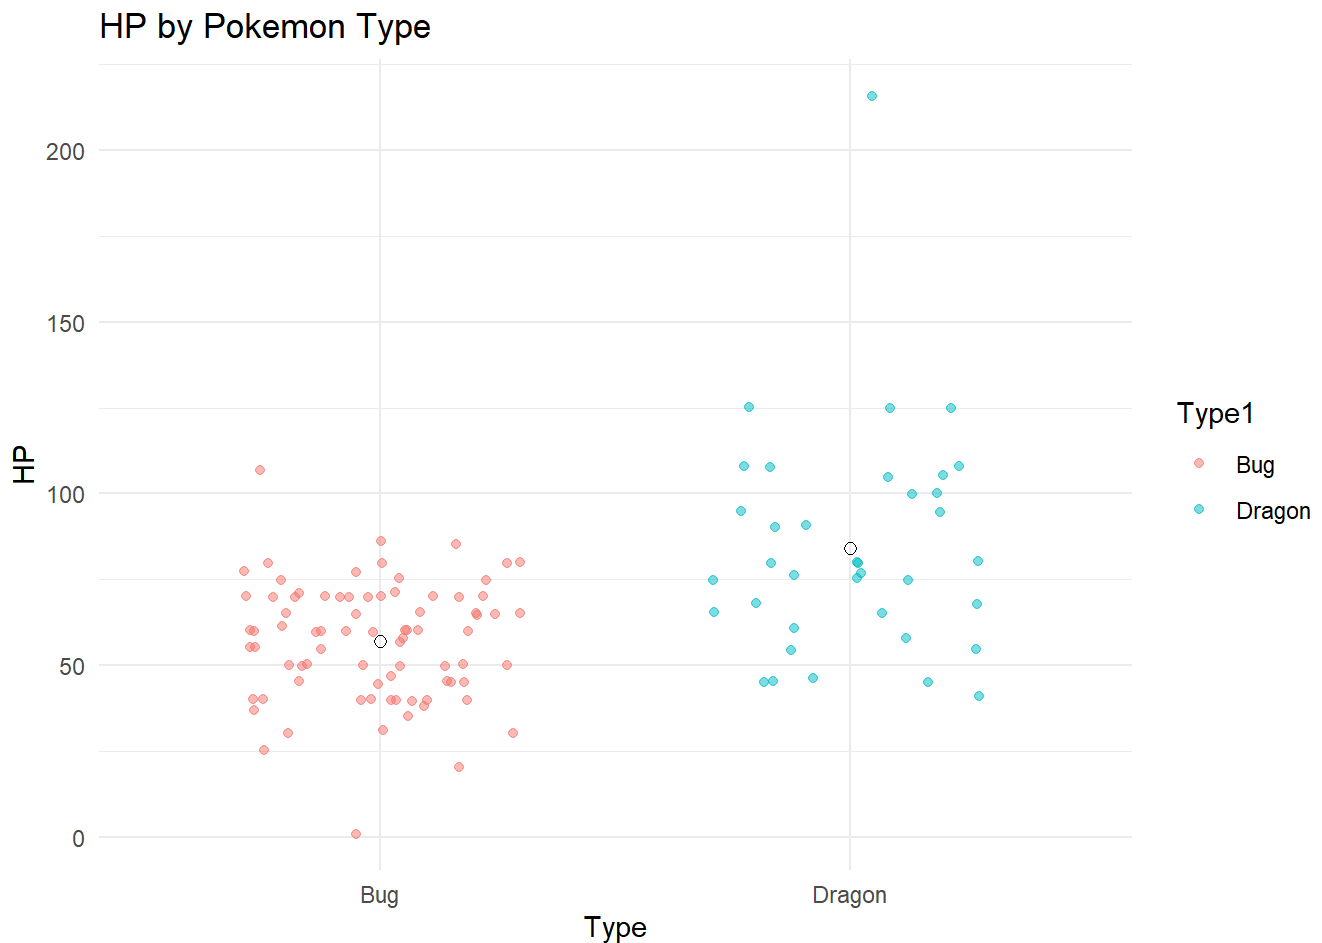
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 23.2 on 115 degrees of freedom

Multiple R-squared: 0.23, Adjusted R-squared: 0.2233

F-statistic: 34.35 on 1 and 115 DF, p-value: 4.481e-08

```
ggplot(pokemon.subsample, aes(x = Type1, y = HP, color = Type1)) +
  geom_jitter(width = 0.3, alpha = 0.5) +
  geom_point(aes(y = fitted(mod.simple)), color = "black", shape = 1, size = 2) +
  theme_minimal() +
  labs(title = "HP by Pokemon Type", x = "Type", y = "HP")
```



I would like to use the second estimator to maximize the difference in HP. As we can see from the two estimators, the second one is the one that features the highest absolute difference. Furthermore, we see that the groups are not balanced, we seem to have more bug pokemon than dragon pokemon. Therefore, we would want to pick the estimator that focuses on the largest subgroup to maximize the difference in HP, which in this case, is the second one.

Exercise #2

```
library(gmm)
```

Warning: package 'gmm' was built under R version 4.3.3

Loading required package: sandwich

Warning: package 'sandwich' was built under R version 4.3.3

```
data("nsw")
```

```
head(nsw)
```

	treat	age	ed	black	hisp	married	nodeg	re75	re78
1	0	23	10	1	0	0	1	0.000	0.00

2	0	26	12	0	0	0	0	0.000	12383.68
3	0	22	9	1	0	0	1	0.000	0.00
4	0	34	9	1	0	0	1	4368.413	14051.16
5	0	18	9	1	0	0	1	0.000	10740.08
6	0	45	11	1	0	0	1	0.000	11796.47

```
summary(nsw)
```

treat	age	ed	black
Min. :0.0000	Min. :17.00	Min. : 3.00	Min. :0.0000
1st Qu.:0.0000	1st Qu.:19.00	1st Qu.: 9.00	1st Qu.:1.0000
Median :0.0000	Median :23.00	Median :10.00	Median :1.0000
Mean :0.4114	Mean :24.52	Mean :10.27	Mean :0.8006
3rd Qu.:1.0000	3rd Qu.:27.00	3rd Qu.:11.00	3rd Qu.:1.0000
Max. :1.0000	Max. :55.00	Max. :16.00	Max. :1.0000

hisp	married	nodeg	re75
Min. :0.0000	Min. :0.000	Min. :0.0000	Min. : 0.0
1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:1.0000	1st Qu.: 0.0
Median :0.0000	Median :0.000	Median :1.0000	Median : 936.3
Mean :0.1053	Mean :0.162	Mean :0.7798	Mean : 3042.9
3rd Qu.:0.0000	3rd Qu.:0.000	3rd Qu.:1.0000	3rd Qu.: 3993.2
Max. :1.0000	Max. :1.000	Max. :1.0000	Max. :37431.7

re78
Min. : 0
1st Qu.: 0
Median : 3952
Mean : 5455
3rd Qu.: 8772
Max. :60308

(a) Average Treatment Effect (ATE)

```
treat <- nsw %>%
  filter(treat == 1)

untreat <- nsw %>%
  filter(treat == 0)

mean(treat$re78) - mean(untreat$re78)
```

```
[1] 886.3037
```

This is the raw difference in mean earnings between those that were treated and those that were untreated across all groups.

(b) Bootstrapped 95% CI (ATE)

```

set.seed(123)

n <- 10000
ate_boot <- rep(NA, n)

for(i in 1:n){
  sample <- nsw[sample(1:nrow(nsw), nrow(nsw), replace = T),]

  treat <- sample %>%
    filter(treat == 1)

  untreat <- sample %>%
    filter(treat == 0)

  ate_boot[i] <- mean(treat$re78) - mean(untreat$re78)
}

```

```

lower.bound <- quantile(ate_boot, 0.025)
upper.bound <- quantile(ate_boot, 0.975)

print(paste0("The Bootstrapped 95% CI is {", lower.bound, ", ", upper.bound, "}"))

```

```
[1] "The Bootstrapped 95% CI is {-56.7367156251803, 1857.96121080787}"
```

If we were to repeat this experiment under the same conditions with a sufficiently large sample size, the average treatment effect would fall somewhere within the estimated interval for about 95 out of 100 trials. In this experiment, expected average treatment effect for our estimator is predicted to be somewhere in between -56.74 and 1857.96.

(c) Least Squares Regression

```

model <- lm(re78 ~ treat, data = nsw)
summary(model)

```

Call:

```
lm(formula = re78 ~ treat, data = nsw)
```

Residuals:

Min	1Q	Median	3Q	Max
-5976	-5090	-1519	3361	54332

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5090.0	302.8	16.811	<2e-16 ***
treat	886.3	472.1	1.877	0.0609 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 6242 on 720 degrees of freedom

Multiple R-squared: 0.004872, Adjusted R-squared: 0.003489

F-statistic: 3.525 on 1 and 720 DF, p-value: 0.06086

```
predicted_re78 <- data.frame(re78 = predict(model), treat = nsw$treat)
```

```
ggplot(nsw, aes(x = treat, y = re78)) +
  geom_point(aes(x = treat, y = re78, color = factor(treat)),
    alpha = 0.3) +
  geom_line(data = predicted_re78,
    aes(x = treat, y = re78), color = 'darkorange', lwd = 1) +
  labs(
    title = "Treatment vs. 1978 Earnings",
    x = "Treatment",
    y = "1978 Earnings",
    color = "Treatment") +
  theme_minimal()
```



The intercept describes the average baseline salary for both treatment and control groups at the start of the study. Finally, the treat coefficient represents the average difference in salary between treated and

untreated individuals. In other words, the average treatment effect. Here, the average treatment effect is estimated to be ~ \$886.

(d) Discussion

I would say that adding certain factors would definitely change the average treatment effect. For instance, the variable of marriage will add an additional source of income. Respondents would likely submit a higher value for the household salary if they were living with their spouse. Therefore, it's important to consider covariates when estimating the average treatment effect.

```
model <- lm(re78 ~ treat + married, data = nsw)
summary(model)
```

Call:

```
lm(formula = re78 ~ treat + married, data = nsw)
```

Residuals:

Min	1Q	Median	3Q	Max
-6510	-4989	-1511	3362	54440

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4989.0	318.7	15.656	<2e-16 ***
treat	879.4	472.1	1.863	0.0629 .
married	641.2	630.5	1.017	0.3095

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6242 on 719 degrees of freedom

Multiple R-squared: 0.006301, Adjusted R-squared: 0.003537

F-statistic: 2.28 on 2 and 719 DF, p-value: 0.1031

(e) Difference in Means (CATE)

```
df <- nsw %>%
  group_by(nodeg) %>%
  summarise(
    N_Total = n(),
    N_Treated = sum(treat == 1),
    N_Control = sum(treat == 0),
    Mean_Treated = mean(re78[treat == 1]),
    Mean_Control = mean(re78[treat == 0]),
    CATE = Mean_Treated - Mean_Control
  ) %>%
  ungroup()
```

```
df
```

```
# A tibble: 2 × 7
```

	nodeg	N_Total	N_Treated	N_Control	Mean_Treated	Mean_Control	CATE
	<dbl>	<int>	<int>	<int>	<dbl>	<dbl>	<dbl>
1	0	159	80	79	6977.	5919.	1059.
2	1	563	217	346	5607.	4901.	706.

These estimates are not causally identified because we are not conditioning treatment on whether an individual has a high school diploma or not.

(f) Standardized CATE

```
df <- df %>%
  mutate(CATE_standard = (N_Total/sum(N_Total))*CATE)

sum(df$CATE_standard)
```

```
[1] 784.0365
```

This value is less than my estimated ATE. This is likely because by conditioning on high school diploma you remove potential discrepancies by education. Therefore, the difference in the treatment is not so stark.

(g) Least Squares Regression

```
model <- lm(re78 ~ treat + nodeg + treat*nodeg, data = nsw)
summary(model)
```

Call:

```
lm(formula = re78 ~ treat + nodeg + treat * nodeg, data = nsw)
```

Residuals:

Min	1Q	Median	3Q	Max
-6977	-4901	-1405	3287	54701

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5918.6	701.0	8.443	<2e-16 ***
treat	1058.7	988.3	1.071	0.284
nodeg	-1017.7	777.0	-1.310	0.191
treat:nodeg	-352.2	1126.0	-0.313	0.755

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6231 on 718 degrees of freedom

Multiple R-squared: 0.01113, Adjusted R-squared: 0.006994

F-statistic: 2.693 on 3 and 718 DF, p-value: 0.04521

Using this least squares regression, the intercept gives the mean earnings for the control group, otherwise known as the baseline. Additionally, the treat coefficient specifies how much change there is in earnings for each unit of treatment. In this case, this represented the CATE for those without a high school diploma. Additionally, the nodeg coefficient represents the change in earnings for each unit of nodeg. In other words, the earnings decrease by ~ \$1017 on average for those with a high school diploma. Finally, the interaction coefficient, represents the relationship between treatment and nodeg, where those that have a high school diploma are less likely to receive treatment. The difference between the treat coefficient and the interaction variable represents the CATE for those without a high school diploma.

(h) Bootstrapped Estimated Sampling Distribution

```
boot_fn <- function(data, indices) {
  d <- data[indices, ]
  fit <- lm(re78 ~ treat + nodeg + treat*nodeg, data = d)
  return(coef(fit))
}
```

```
set.seed(123)
boot_results <- boot(data = nsw, statistic = boot_fn, R = 10000)
```

```
boot_results
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = nsw, statistic = boot_fn, R = 10000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	5918.6075	-8.167754	691.9999
t2*	1058.7049	17.062629	1106.7927
t3*	-1017.7389	8.547283	755.0760
t4*	-352.2391	-15.372697	1230.9072

```
boot.ci(boot_results, type = "perc", index = 2)
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 10000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_results, type = "perc", index = 2)
```

Intervals :

Level	Percentile
-------	------------

95%	(-1003, 3312)
-----	----------------

Calculations and Intervals on Original Scale

```
boot.ci(boot_results, type = "perc", index = 4)
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 10000 bootstrap replicates

CALL :

```
boot.ci(boot.out = boot_results, type = "perc", index = 4)
```

Intervals :

Level	Percentile
-------	------------

95%	(-2844.3, 1977.1)
-----	--------------------

Calculations and Intervals on Original Scale

If we were to repeat this experiment under the same conditions with a sufficiently large sample size, the conditional average treatment effect would fall somewhere within the estimated interval for about 95 out of 100 trials. In this experiment, expected conditional average treatment effect for those without a high school diploma is predicted to be somewhere in between -1003 and 3312. Also, the expected conditional average treatment effect for this experiment for those with a high school diploma is expected to fall somewhere between -1841 and 1335 (if you were to subtract the values from both CI above).